

**CS5250 – Advanced Operating Systems
AY2018/2019 Semester 2**

Assignment 4

**Deadline: Friday, 19 April 2019 • 11.59pm
This is the last assignment for this course.**

Section A (25 marks)

1. Objectives

1. Get a deeper understanding on process scheduling.

2. Rules

1. It is fine to ask for “reasonable” amount of help from others, but ensure that you do all the tasks on your own and write the report on your own. The University’s policy on plagiarism applies here and any breaches will be dealt with severely.
2. For this assignment, you are asked to finish three tasks, and write a report (check assignment section for more details).
3. Generate your report as a pdf file, name it as “*Name (Student Number) Assignment4.pdf*” and upload your report in the IVLE folder Assignment-4 of CS5250.
4. The deadline is 19 April 2019. Late assignments lose 4 marks per day.
5. For any question, contact the teaching assistant, Mr. Vanchinathan Venkataramani, a0048048@u.nus.edu or Mr. Chen Cheng e0205030@u.nus.edu.

3. Assignment

In this assignment, we will study the effect of process scheduling scheme on system’s performance. In particular, the average waiting time on each process. First, we will build a simulator for scheduling tasks. Next, we explore/compare different scheduling schemes to draw conclusion on how effective they are.

Task 1: 12 Marks:

Complete the simulation for

1. Round Robin (RR) (4 Marks)
2. Shortest remaining time first (SRTF) with given cpu burst time (4 Marks)
3. **Shortest job first (SJF)** with future prediction (4 Marks)

Task 2: 13 Marks:

You are to create your own simulator and experiment with different input/configuration and answer the questions.

Details:

Task 1:

A sample simulator with FCFS schemes was completed for your reference (see link below). You can choose to work on this file, or write your own simulator in your favorite language.

Please make sure input and output format are still the same as in the sample program and as described in this document.

Input:

For the simulation, we will assume that all tasks are CPU bound. The actual cpu time required to complete the tasks are given in the input.txt file.

The **input.txt** file provides the list of processes which will be run on our simulator.

The file contains several lines, each line represents 1 process with 3 integer values:
<process_id arrival_time burst_time>

For example:

1 0 2

means that the process with id = 1 arrives at the scheduler at time $t = 0$, and completes execution after 2 time units. You can assume that the time unit is in milliseconds (ms).

Time $t = 0$ is when the simulation starts, there will always be a process at time $t=0$;

The file **input.txt** lists all the processes in the increasing order of their arrival time. Assume only one process arrives at a time slice (if process 0 arrives at time 0, process 1 must arrive at time > 0 (any number > 1)).

Output:

Our scheduler will decide when to switch from the current process to another process. This switching depends on the scheduling scheme used. To record the result of the simulator, we will output the time stamp at each process switching event (only when the new process is different from the current process), i.e. (time, process id).

Please assume that context switching overhead is 0, and the time starts at 0ms.

Using FCFS for this input:

0 0 4
1 2 3

Will produce this output:

(0, 0)
(4, 1)
Average waiting time 1.0

In the above example, the input means: process id = 0 arrives at time $t = 0$, will take 4 time units to complete. Process id = 1 arrives at time $t=2$, will take 3 time units to complete.

Using FCFS, the scheduler will switch to process 0 at time $t=0 \Rightarrow$ line 1 (0,0) in the output. After process 0 complete then process 1 will be switched in by the scheduler at time $t=4 \Rightarrow$ line 2: (4,1).

The average waiting time is 1.0 because process 0 didn't have to wait, process 1 has to wait 2 time units (from the time it arrives $t=2$ to the time it starts $t=4$). On average, the waiting time is $(0 + 2)/2 = 1.0$

The given code has the sample for FCFS, your task is to complete the other 3 scheduling schemes:

1. Round robin with time quantum Q . Q is the parameter for your function,
2. SRTF with the given CPU burst time
3. SJF (non-preemptive) with the future-prediction formula in Lecture 7, slide 48-50. Using $\alpha = 0.5$, initial guess = 5 for all processes.

For simplicity, we will perform non-preemptive SJF with future prediction. If the current process is scheduled, we wait until it finishes before scheduling another process. We need to keep the record of previous burst time for each process, then predict its current burst time when it arrives at the scheduler. The prediction is per-process and based on their execution history with the initial guess = 5 time units.

Commit the code for the above schemes on your Github and submit the link to the report.

Please also report the schedule output and the average waiting time for each of the four scheduling algorithms, i.e. FCFS and Tasks 1.1-1.3. Explain your observations.

Task 2:

1. Given the below input:

```
0 0 9
1 1 8
2 2 2
3 5 2
3 30 5
1 31 2
2 32 6
0 38 8
2 60 7
0 62 2
1 65 3
3 66 8
1 90 10
0 95 10
2 98 9
3 99 8
```

Test with all the three implemented scheduling schemes (Task 1.1-1.3) and compare them. Which one gives the least average_waiting_time?

Try to adjust the time quantum Q for RR and α value for SJF, give the optimal value of Q and α to minimize average_waiting_time for the above particular input. Include the output of your tests in the report.(4 marks)

2. Which of the evaluated schemes in this assignment (Task 1.1-1.3, 2.1) generates the optimal schedule (gives minimum average waiting time) for a system with (4 marks):
 - a) All short processes

- b) Very short and very long processes interleave each other with unpredictable pattern.
3. State a different scheduling mechanism that has not been mentioned in the lecture notes and explain the intuition behind the using this scheme. (2 marks)
4. Assume your system has N CPU cores, and each process only requires burst time on 1 core. Will it make the scheduler more complicated? Suggest how to extend the current scheduler to multi-processor system. (3 marks)

Sample Code for Task 1 & 2:

Sample code is provided here:

<https://github.com/vanchiramani/cs5250-assignment4.git>

Compatible with python 2 & 3. To run the simulation, change the input.txt file, then simply type
python simulator.py

Section B (10 marks)

Consider the following compare-and-exchange atomic instruction (abstracted as a function – in other words, assume that the following function is atomic):

int cas(int *ptr, int oldval, int newval);

If the integer pointed to by **ptr** is equal to **oldval**, then the content pointed to by **ptr** will be replaced with **newval**, and a '1' will be returned by **cas**. Otherwise, a '0' is returned. Using this function/instruction, show how a spinlock can be implemented by giving the C code for the **spin_lock()** and **spin_unlock()** functions. Write down any assumptions that you make.

- End of Assignment -